



# Vidyavardhini's College of Engineering & Technology

## Department of Artificial Intelligence and Data Science

AY: 2025-26

<b>Class:</b>	BE	<b>Semester:</b>	VII
<b>Course Code:</b>	CSDOL7011	<b>Course Name:</b>	Natural Language Processing

<b>Name of Student:</b>	Konisha Jayesh Thakare
<b>Roll No. :</b>	71
<b>Experiment No.:</b>	10
<b>Title of the Experiment:</b>	Design and Development of a Real-World NLP Application
<b>Date of Performance:</b>	16.09.2025
<b>Date of Submission:</b>	23.09.2025

### Evaluation

Performance Indicator	Max. Marks	Marks Obtained
Performance	5	
Understanding	5	
Journal work and timely submission	10	
Total	20	

Performance Indicator	Exceed Expectations (EE)	Meet Expectations (ME)	Below Expectations(BE)
Performance	4-5	2-3	1
Understanding	4-5	2-3	1
Journal work and timely submission	8-10	5-8	1-4

### Checked by

**Name of Faculty :** Dr. Tatwadarshi P. Nagarhalli

**Signature :**

**Date :**



**Aim:** To design and develop a real-world NLP application, implementing NLP techniques learned in previous experiments, and to evaluate its performance and usability.

### **Dataset Details:**

**Source:** [IMDB Dataset of 50K Movie Reviews](#)

**Size:** 50,000 movie reviews

### **Columns:**

- review: Text of the movie review
- sentiment: Label (positive or negative)

**Train/Test Split:** 80% training (40,000 reviews), 20% testing (10,000 reviews)

### **Theory:**

Sentiment analysis is a classic NLP task to determine the sentiment expressed in text. In this experiment:

1. Text Preprocessing removes noise from raw reviews: lowercasing, removing punctuation, stopwords removal, and stemming using the Porter Stemmer.
2. Feature Extraction converts text into numerical vectors using TF-IDF, capturing importance of words relative to the document and corpus.
3. Model Training uses Multinomial Naïve Bayes, suitable for text classification with count-based features.
4. Evaluation is done using metrics:
  - Accuracy: Overall correctness of predictions
  - Precision, Recall, F1-score: Performance per class
  - Confusion Matrix: Correct vs. misclassified examples

### **Procedure:**

1. Load IMDB dataset and map positive  $\rightarrow$  1, negative  $\rightarrow$  0.
2. Preprocess reviews:
  - Convert to lowercase
  - Remove punctuation



- Tokenize
  - Remove stopwords
  - Apply Porter Stemming
3. Convert processed reviews into TF-IDF vectors.
  4. Split dataset into training (80%) and testing (20%) sets.
  5. Train Multinomial Naïve Bayes classifier on the training set.
  6. Predict sentiment on the test set.
  7. Evaluate model performance using accuracy, classification report, and confusion matrix.
  8. Build a simple interface function to predict sentiment of new user input.

### Implementation:

```
import nltk
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

nltk.download('stopwords')
stop_words = set(stopwords.words('english'))
ps = PorterStemmer()

data = pd.read_csv("IMDB_Reviews.csv")
data['sentiment'] = data['sentiment'].map({'positive':1, 'negative':0})

def preprocess(text):
    text = text.lower()
    text = text.translate(str.maketrans('', '', string.punctuation))
    tokens = nltk.word_tokenize(text)
    tokens = [ps.stem(word) for word in tokens if word not in stop_words]
    return " ".join(tokens)

data['clean_text'] = data['review'].apply(preprocess)
tfidf = TfidfVectorizer()
X = tfidf.fit_transform(data['clean_text'])
y = data['sentiment']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```



# Vidyavardhini's College of Engineering & Technology

## Department of Artificial Intelligence and Data Science

```
model = MultinomialNB()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

def predict_sentiment(text):
    text = preprocess(text)
    vector = tfidf.transform([text])
    pred = model.predict(vector)
    return "Positive" if pred[0] == 1 else "Negative"

user_input = "I loved the movie, it was fantastic!"
print("Sentiment Prediction:", predict_sentiment(user_input))
```

### Results:

```
PS C:\Users\Konisha Thakare\OneDrive\Desktop\Python> python imdb_sentiment_classifier.py
[nltk_data] Downloading package stopwords to C:\Users\Konisha
[nltk_data]   Thakare\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
Accuracy: 0.8641

Classification Report:
              precision    recall  f1-score   support

     0       0.85         0.88         0.86         4961
     1       0.88         0.85         0.86         5039

 accuracy          0.86
 macro avg         0.86
weighted avg         0.86

Sentiment Prediction: Positive
```

### Example User Input:

```
"I loved the movie, it was fantastic!" → Predicted Sentiment: Positive
```

### Interpretation:

- High accuracy (~86%) demonstrates the model can reliably distinguish positive and negative reviews.
- Balanced precision and recall indicate consistent performance across classes.
- The application can predict sentiment for unseen movie reviews effectively.



### **Conclusion:**

This experiment successfully demonstrated the design and development of a real-world NLP application for sentiment analysis using the IMDB 50K movie reviews dataset. By combining preprocessing, TF-IDF feature extraction, and Multinomial Naïve Bayes classification, the model achieved 86.4% accuracy on the test set. The application can predict the sentiment of new reviews in real-time, demonstrating the practical utility of NLP techniques for real-world scenarios like movie review analysis. Proper experimental methodology, including data preprocessing, training/testing split, and performance evaluation, ensures reliable and reproducible results.